

eda_data_cleaning

December 18, 2025

1 Introduction

Predicting loan approval is an important problem in consumer facing finance. Bank lenders can receive thousands to hundreds of thousands of applications from individuals with a wide variety of financial variables, leading to the necessity for automation of these loan approval processes. Approving loans which end up defaulting can increase financial risk, while being too conservative in lending loans can reduce business opportunity and leave valuable profits to waste. As a way to reduce these errors in the automated loan approval process, organizations rely on data driven methods and algorithms to get a better understanding of loan approval factors.

In this project, we will explore a synthetic loan approval dataset and identify key drivers behind approval decisions, allowing us to build a predictive machine learning model to accurately determine a loan applications approval status.

2 Adding Directory For Plots

```
[1]: import os  
os.makedirs("plots", exist_ok=True)
```

3 Data Description

The dataset contains information on thousands of past loan applications, each row representing an individual applicant, with a mix of variables that may be related to loan approval status.

This data contains 20 features (1 id, 1 target), a distribution of 55/45 in target status, and includes real world financial approval logic (DTI ratios, Filed Defaults, etc...)

This dataset provides a healthy combination of demographic, financial, and credit related information, making it suitable for understanding patterns in loan approvals. Before modeling, we will clean the dataset, address missing values, check for outliers, and explore some potential relationships between key features and loan outcomes.

4 Data Processing and Exploration

```
[2]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
import plotly.express as px
```

```
[3]: df = pd.read_csv("data/Loan_approval_data_2025.csv")  
df.head(5)
```

```
[3]:   customer_id  age occupation_status  years_employed  annual_income  \  
0    CUST100000  40      Employed          17.2        25579  
1    CUST100001  33      Employed          7.3         43087  
2    CUST100002  42      Student           1.1        20840  
3    CUST100003  53      Student           0.5        29147  
4    CUST100004  32      Employed          12.5       63657  
  
      credit_score  credit_history_years  savings_assets  current_debt  \  
0            692                  5.3        895       10820  
1            627                  3.5        169       16550  
2            689                  8.4        17        7852  
3            692                  9.8       1480      11603  
4            630                  7.2        209       12424  
  
      defaults_on_file  delinquencies_last_2yrs  derogatory_marks  product_type  \  
0                  0                      0                0  Credit Card  
1                  0                      1                0  Personal Loan  
2                  0                      0                0  Credit Card  
3                  0                      1                0  Credit Card  
4                  0                      0                0  Personal Loan  
  
      loan_intent  loan_amount  interest_rate  debt_to_income_ratio  \  
0     Business        600      17.02        0.423  
1  Home Improvement     53300      14.10        0.384  
2  Debt Consolidation      2100      18.33        0.377  
3     Business        2900      18.74        0.398  
4     Education       99600      13.92        0.195  
  
      loan_to_income_ratio  payment_to_income_ratio  loan_status  
0              0.023                  0.008          1  
1              1.237                  0.412          0  
2              0.101                  0.034          1  
3              0.099                  0.033          1  
4              1.565                  0.522          1
```

```
[4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customer_id      50000 non-null   object  
 1   age              50000 non-null   int64  
 2   occupation_status 50000 non-null   object  
 3   years_employed    50000 non-null   float64 
 4   annual_income     50000 non-null   int64  
 5   credit_score      50000 non-null   int64  
 6   credit_history_years 50000 non-null   float64 
 7   savings_assets    50000 non-null   int64  
 8   current_debt      50000 non-null   int64  
 9   defaults_on_file  50000 non-null   int64  
 10  delinquencies_last_2yrs 50000 non-null   int64  
 11  derogatory_marks  50000 non-null   int64  
 12  product_type      50000 non-null   object  
 13  loan_intent       50000 non-null   object  
 14  loan_amount       50000 non-null   int64  
 15  interest_rate     50000 non-null   float64 
 16  debt_to_income_ratio 50000 non-null   float64 
 17  loan_to_income_ratio 50000 non-null   float64 
 18  payment_to_income_ratio 50000 non-null   float64 
 19  loan_status       50000 non-null   int64  
dtypes: float64(6), int64(10), object(4)
memory usage: 7.6+ MB

```

[5]: df = df.drop('customer_id', axis=1)

4.1 Basic EDA

4.1.1 Target Variable Balance

[6]: df["loan_status"].value_counts().reset_index()

```

[6]:   loan_status  count
 0            1  27523
 1            0  22477

```

```

[7]: # Get value counts
loan_counts = df["loan_status"].value_counts().sort_index()

# Create the figure
fig, ax = plt.subplots(figsize=(7, 5))

# Create bar plot

```

```


ax.bar(loan_counts.index, loan_counts.values,
       color=['#EF553B', '#636EFA'], # Red and blue to match Plotly colors
       width=0.6)

# Customize
ax.set_xlabel('loan_status', fontsize=11)
ax.set_ylabel('count', fontsize=11)
ax.set_title('Loan Status Distribution', fontsize=12)
ax.set_xticks([0, 1])

# Optional: add gridlines for easier reading
ax.grid(axis='y', alpha=0.3, linestyle='--')
ax.set_axisbelow(True)

plt.tight_layout()
plt.savefig("plots/loan_status_distribution.png")
plt.show()

```




4.1.2 Univariate EDA for numerical columns

```
[8]: num_col = ["annual_income", "loan_amount", "credit_score",
              "debt_to_income_ratio", "loan_to_income_ratio",
              "payment_to_income_ratio", "credit_history_years"]

n = len(num_col)
rows = (n + 1) // 2
plt.figure(figsize=(12, rows * 4))

for i, col in enumerate(num_col, 1):
    plt.subplot(rows, 2, i)
    sns.histplot(df[col])
    plt.title(f"Histogram of {col}")
    plt.savefig(f"plots/{col}_histogram.png")

plt.tight_layout()
plt.show()
```



4.1.3 Correlation Heatmap

```
[9]: numeric_df = df.select_dtypes(include="number")


corr = numeric_df.corr()

fig, ax = plt.subplots(figsize=(12, 10))

# Create heatmap
sns.heatmap(corr,
            annot=True, # Show correlation values
            fmt='.3f', # Format to 2 decimal places
            cmap='RdBu_r', # Red-Blue reversed colormap
            center=0, # Center colormap at 0
            square=True, # Make cells square
            linewidths=0.5, # Add gridlines
            cbar_kws={"shrink": 0.8},
            ax=ax)

ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right', fontsize=9)
ax.set_yticklabels(ax.get_yticklabels(), rotation=0, fontsize=9)
ax.set_title('Correlation Heatmap', fontsize=14, pad=20)

plt.tight_layout()
plt.savefig("plots/correlation_heatmap.png", dpi=150, bbox_inches='tight')
plt.show()
```



4.1.4 Bivariate Analysis

```
[10]: import matplotlib.pyplot as plt

categorical_cols = ['occupation_status', 'product_type', 'loan_intent']

for col in categorical_cols:
    # Get value counts for each loan status
    counts_0 = df[df['loan_status'] == 0][col].value_counts().sort_index()
    counts_1 = df[df['loan_status'] == 1][col].value_counts().sort_index()

    # Initializing figure to plot on
    fig, ax = plt.subplots(figsize=(8, 5))

    # Bar position configuration
    categories = sorted(df[col].unique())
```

```


x = range(len(categories))
width = 0.35



# Bars per group
ax.barh([i - width/2 for i in x],
        [counts_0.get(cat, 0) for cat in categories],
        width, label='loan_status=0', color="#EF553B")
ax.barh([i + width/2 for i in x],
        [counts_1.get(cat, 0) for cat in categories],
        width, label='loan_status=1', color="#636EFA")

# Axis configuration
ax.set_yticks(x)
ax.set_yticklabels(categories)
ax.set_xlabel('count')
ax.set_ylabel(col)
ax.set_title(f"{col} vs Loan Status")
ax.legend()
ax.grid(axis='x', alpha=0.3, linestyle='--')
ax.set_axisbelow(True)

plt.tight_layout()
plt.savefig(f"plots/{col}_vs_loan_status.png", dpi=150, bbox_inches='tight')
plt.show()

```






```
[11]: n = len(num_col)
rows = (n + 1) // 2
```

```
plt.figure(figsize=(12, rows * 4))

for i, col in enumerate(num_col, 1):
    plt.subplot(rows, 2, i)
    sns.boxplot(x=df['loan_status'], y=df[col])
    plt.title(f"Boxplot of {col} by Loan Status")
    plt.savefig(f"plots/{col}_boxplot.png")

plt.tight_layout()
plt.show()
```



```
[12]: average_ratios = df.groupby("loan_status")[num_col].median().T
average_ratios
```

```
[12]: loan_status          0          1
      annual_income     37127.000  45708.000
      loan_amount       26900.000  25400.000
      credit_score        603.000   671.000
      debt_to_income_ratio  0.329    0.232
      loan_to_income_ratio  0.726    0.539
      payment_to_income_ratio  0.242    0.180
      credit_history_years   3.600    8.600
```

4.2 Insight from EDA

Some details and insights that may be useful or interesting gathered from this exploratory data analysis.

- Most approved lendee's are on average older, longer employed, and higher credit score
- Most approved lendee's have lower debt
- Most approved lendee's have a better DTI, LTI, and PTI ratio (lower = better)
- Credit card loans on average have a higher rate of denial
- Credit history on average approved lendee's is 2x higher than those denied

Here we will convert our cleaned data into a csv for model training and evaluation:

```
[13]: df.to_csv('data/cleaned_data.csv', index=False)
```